

# Firefighting a Symfony & Elasticsearch app with Blackfire

Symfony User Group - Zürich

2024 - 03-04

Hello!

# Emanuele Panzeri | thePanz

Software Engineer at Liip - Zurich

Elastica PHP Client maintainer

[github.com/thePanz](https://github.com/thePanz)

[@thepanz@phpc.social](mailto:@thepanz@phpc.social)



# Agenda



- Context 🐙
- Application is on fire 🔥
- Firefighting sessions 🚒
- Lessons learned 📖

# Application Context

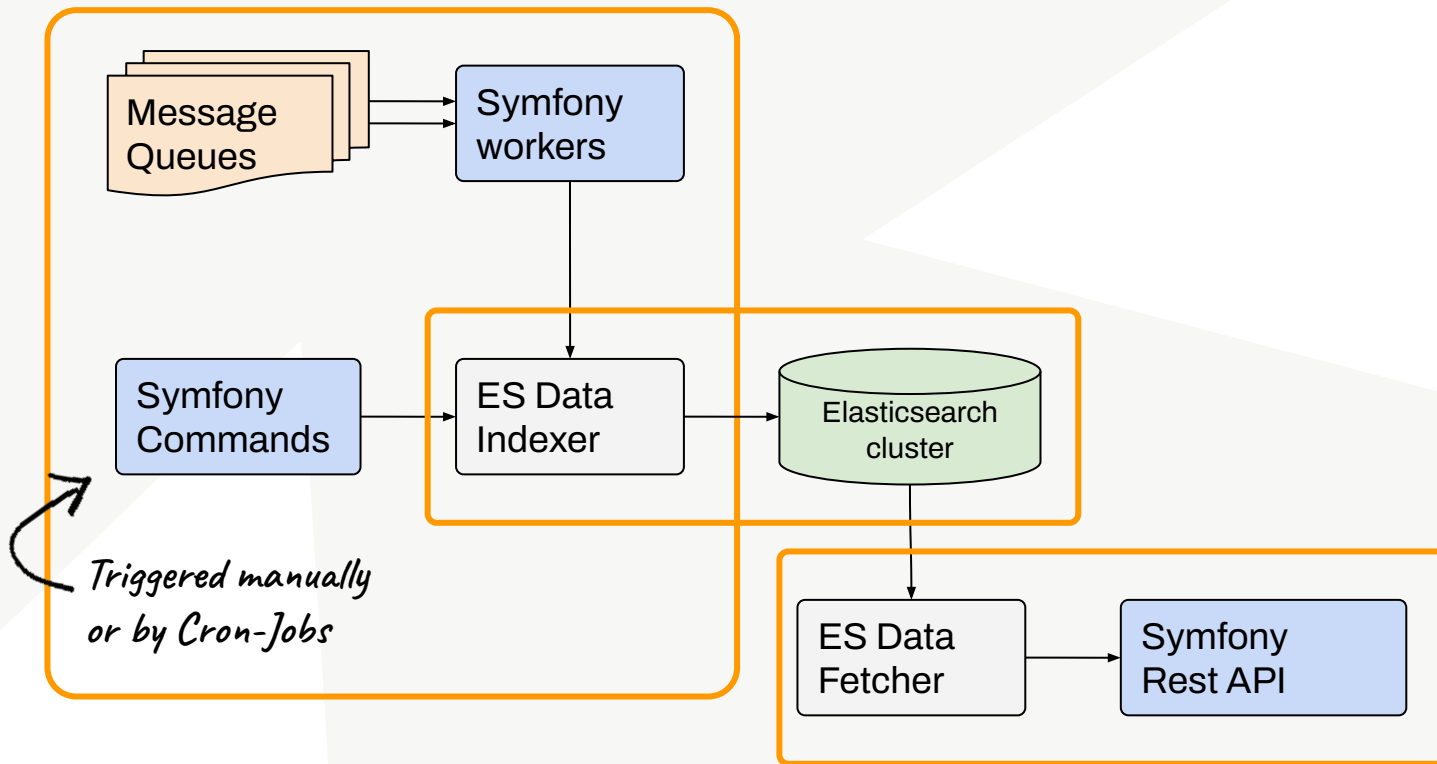
***“We built a data kraken”***

**— dbu @ SymfonyCon 2022**

# Product API built on Symfony

- Uses multiple SF components  
*Messenger, HttpClient, Process, Notifier, Validator, ..*
- Aggregates data from 30+ sources with multiple formats  
*JSON, XML, CSV accessed from REST, XMLRPC, ...*  
*And some other weird ones too*  

- Consolidates (and clears) source information  
*Database for local caching, Elasticsearch for aggregations and search*
- Exposes data via JSON REST APIs  
*Serving ~35M requests per day (mobile, cashier registries, ...)*  
*Most read-only APIs*  


# Data Workflow: acquisition and exposing



# Data Workflow: ingestion

- Data import (into MySQL, Redis)
  - Incoming RabbitMQ messages for data updates
  - Data importing tasks run via Cron-jobs
  - Data changes dispatch (internal) messages for indexing

- Data indexing
  - Entities are build by aggregating source data
  - Denormalized objects are stored into Elasticsearch
  - Entity-changed events trigger outgoing messages

*We do partial updates  
of ES objects*



 Elasticsearch is on fire

*Data Too Large Exception*





# Elasticsearch is on fire

We started noticing that mostly for **Products**-related areas:

- Queues messages piled up
- Product updates were delayed
- API responses returned HTTP-500

In business speaking 

- Customers were not able to purchase products online
- Prices were not updated online and on the store e-labels
- Personal discounts were not available at the cashiers

# Elasticsearch is on fire

Elasticsearch errors started to appear on our logs for:

- Indexing pipelines
- delivering API Responses

```
update: *****_product_de_20221208/_doc/131415 caused  
[parent] Data too large, data for [indices:data/write/bulk[s]] would be [31648082036/29.4gb],  
which is larger than the limit of [31621696716/29.4gb],  
real usage: [31648077680/29.4gb],  
new bytes reserved: [4356/4.2kb],  
[...]
```

Elastica\Exception\Bulk\Response\ActionException

Which translated as:

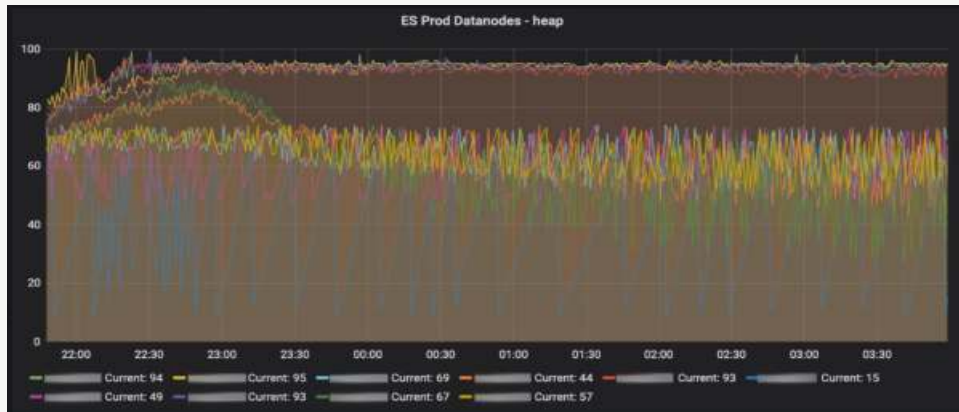
*“The indexing of product 131415  
required more than 29.4GB of memory”* 🤯

*The product itself is 4.2kb  
of JSON data*



# 🔥 Elasticsearch is on fire

On the Elasticsearch cluster, the situation was not good either

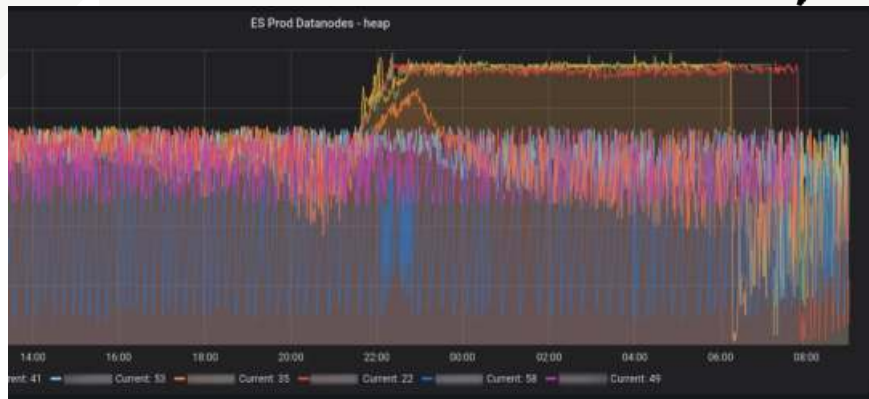


- The cluster was unresponsive
- 3 nodes were blocked, with 95% of memory used  
→ which blocked the entire cluster to process other requests

# 🔥 Elasticsearch is on fire

What happened?

- No previous updates on the ES Cluster
- No changes on the application side
- Restarting the affected nodes solved the issue
  - Until few days later



# ES: Data Too Large?

*“Something else is holding on  
to excessive amounts of memory”* 🙄

<https://discuss.elastic.co/t/what-does-this-error-mean-data-too-large-data-for-transport-request/209345/6> (from 2019)

*“Try increasing the available memory  
of your ES Cluster”*

<https://www.elastic.co/guide/en/elasticsearch/reference/7.14/fix-common-cluster-issues.html#circuit-breaker-errors>


Did not work and the Cluster's JVM was configured according to ES best practices

# 🔥 Firefighting

*Analyze, Log and Profile*



# Emergency measures

- Reduce the number of parallel workers
  - delayed updates to prices still caused issues
- Upgrade ES to latest patch version
  - Not something you want to blindly do in those situations
  -  ES Data not compatible when downgrading a patch version
- Dev-Ops team restarting ES nodes 24/7 was not sustainable

# Next: Analyze and profile

Start a focused the analysis on the Application

- Gather logs and metrics
- Avoid overloading ES with too many requests
- Optimizing the data flow



🔥 Firefighting

*Logs and Deprecations*



# ElasticSearch Logs - I

Our centralized logs kept some of the ES logs:

- Most from high-level logging → nothing helpful
- Our docker instance provided helpful information, tho!

*Deprecated: Do not use the `_id` field for aggregations, sorting, and scripting as it requires to load a lot of data in memory.*

*Does not sound critical*



<https://www.elastic.co/guide/en/elasticsearch/reference/7.17/mapping-id-field.html>

<https://github.com/elastic/elasticsearch/pull/64610>




The `_id` field was our “default” sorting when nothing was provided!

# ElasticSearch Logs - II

## Learnings 💡

- Use a dedicated 'id' field to reduce memory pressure on ES cluster
- Make sure that *\*all\** Logs are collected (deprecations too)
  - ES can collect deprecations in “hidden” indexes in the cluster
- Do not underestimate deprecations!

*Helped to lower the average  
memory usage...*



# 🔥 Firefighting

## *Profile and Optimize With Blackfire*



# App: Precompute - I

Analyzed the impact on ES of our Rest APIs

- Multiple ES queries per “product”
- Additional ES query per aggregation/facet

Examples:

- Additional information computed from ES
- Extend aggregations data for API presentation

*Term and Count are provided by ES, the rest is fetched from the Brand ES index*


```
{
  # Example API response for Products collection
  "facets" : {
    "brand": {
      "name": "Brands",
      "terms": [
        { "term": "PT15", "count": 54, "name": "Best Potato", "slug": "best-potato" },
        { "term": "PZ41", "count": 68, "name": "Pizza Yum!", "slug": "pizza-yum" }
      ]
    }
  },
  "hits" : [ ... ],
}
```

# App: Precompute - II

```
{
  # Raw Product JSON data stored in ES
  "name": "Cheese and Spinach",
  ...
  "brand": {
    "id": "PZ15",
    "name": "Pizza Yum!",
    "slug": "pizza-yum"
    "facet_value": "{\\"id\\":\\"PZ41\\",\\"name\\":\\"Pizza Yum!\\",\\"slug\\":\\"pizza-yum\\"}"
  },
}
```

*ES uses "facet\_value" for aggregation  
providing all data needed for our API*

*~40% less queries on ES*



## Learnings 💡

- Heavily denormalize your data in ES
- Build aggregations on ad-hoc ES fields

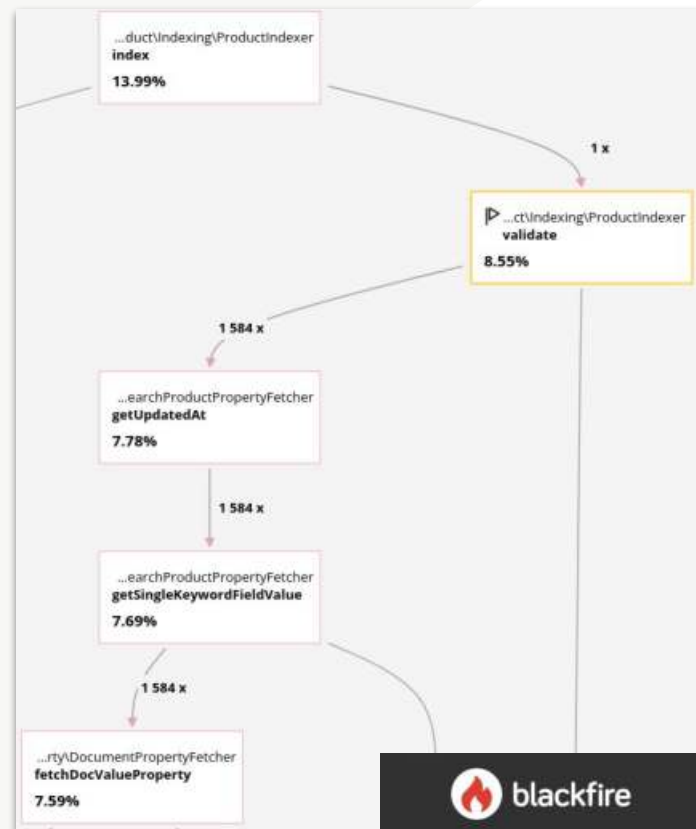
# App: Batch fetch what you need - I

Used **Blackfire** to profile an “mid-average” product indexing (~300 variants)

Why so many calls? *For “legacy” reasons*

After sending data to ES:

- We (re-)fetch the affected products
- Validate that the *updated\_at* property is correct
- Products were fetched 1 by 1 (for all languages)




# App: Batch fetch what you need - II

Improvements on “index” call:

- ~ 1500 less HTTP requests to ES
- ~60% less time (-15s)
- ~80% less IO wait
- ~7% less data transfer

71x	Gone	-24 ms	-60 kB	http://	ch:9200/	product_it_20200914/_search
70x	Gone	-17 ms	-59 kB	http://	ch:9200/	product_de_20200914/_search
68x	Gone	-26 ms	-57 kB	http://	ch:9200/	product_fr_20200914/_search
74x	Gone	-22 ms	-62 kB	http://	ch:9200/	product_de_20200914/_search
72x	Gone	-49 ms	-61 kB	http://	ch:9200/	product_en_20200914/_search
76x	Gone	-18 ms	-64 kB	http://	ch:9200/	product_it_20200914/_search
77x	Gone	-19 ms	-65 kB	http://	ch:9200/	product_en_20200914/_search
76x	Gone	-20 ms	-64 kB	http://	ch:9200/	product_en_20200914/_search
81x	Gone	-20 ms	-68 kB	http://	ch:9200/	product_fr_20200914/_search
78x	Gone	-24 ms	-66 kB	http://	ch:9200/	product_it_20200914/_search





# App: Batch fetch what you need - III

## Learnings

- Batch fetch data from ES
- Directly fetch data from the ES DocValues<sup>2</sup>
  - Do not ask ES to parse the JSON object
- Paginate results without creating memory overhead<sup>1</sup>
  - Use ES “*search\_after*”<sup>2</sup> feature

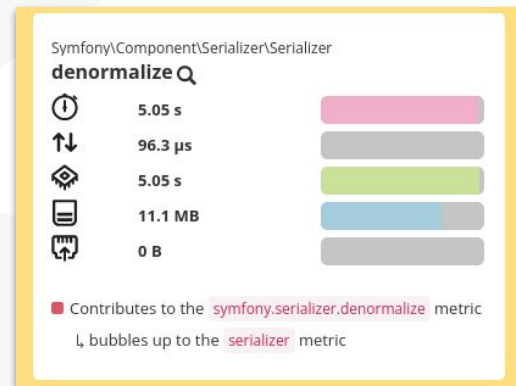
[1] <https://www.elastic.co/guide/en/elasticsearch/reference/current/paginate-search-results.html#search-after>

[2] <https://www.elastic.co/guide/en/elasticsearch/reference/current/doc-values.html>

# App: Do not over (de)Serialize - I

**Blackfire:** profile data building for Product prices  
(product with 5 variants)

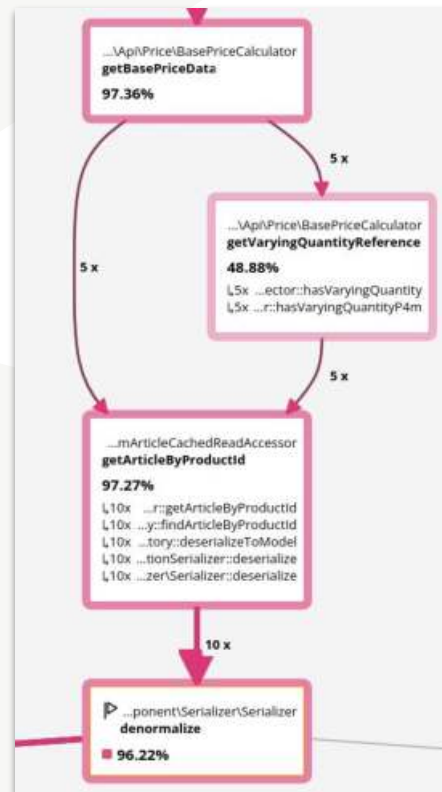
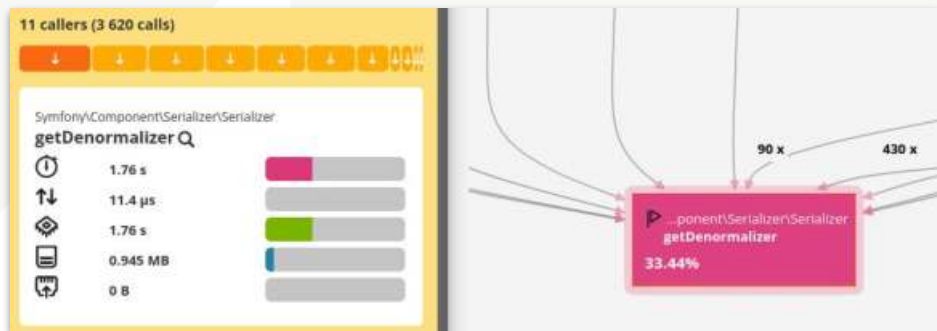
- Product JSON data is stored in MySQL
  - “denormalized” into PHP object with SF Serializer
- We use JanePhp to build PHP Models from OpenAPI
  - Normalizers are generated for schema objects
- A highly nested structure in the JSON creates a bottleneck for *denormalize()*



# App: Do not over (de)Serialize - II

- We know that SF Serializer is fast
- But handling ~300 different models/normalizers might be too much :)

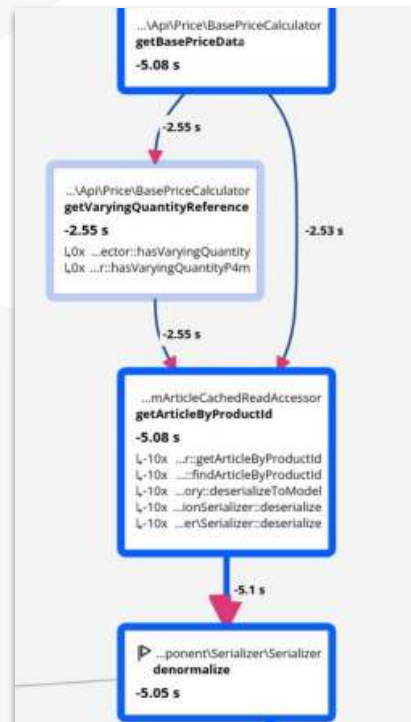
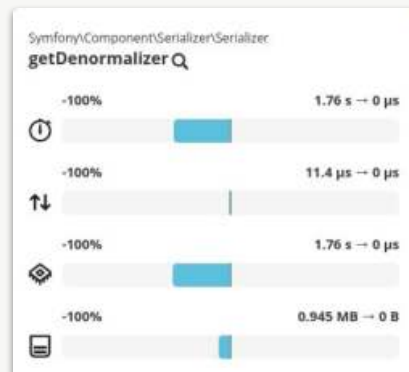
*We are “just” the consumers of this OpenAPI schema  
No chance to get a cleaner structure :(*



# App: Do not over (de)Serialize - III

Improvements:


- Cache deserialized models in Redis
- Use EventDispatcher to prune the cache
- +128 kB (+62.7%) network IO to Redis
- ~97% less execution time (-5s)
- ~70% less memory used (-11MB)



# App: Do not over (de)Serialize - IV

## Learnings 💡

- Trust the SF components, but keep an eye on performances
  - 🙌 from SF 6.4 the ``getDenormalizer()`` call is now cached!
- Try to simplify your schema models (if you can)
- Clear the caches!

*There are only two hard things in Computer Science:  
and off by one errors*  **cache invalidation** *and naming things.*

— Phil Karlton



# Are we on track?

*Heading to the  
right direction*




# App: Performance gains

The changes introduced dramatically improved the application

- Workers were faster
  - Less time to consume “indexing” messages
  - Lower memory footprint
- Less ES requests per worker
- Less ES queries to build API responses

# App: Performance gains, but... 🤔

- More frequent memory spikes on the ES Cluster
- Had to further reduce the max number of product-workers
  - Before: max 50
  - After: max 20





**Shift focus**

***Can we trust the Cloud?***

# Shift focus: Servers and ES Cluster

- Self hosted by another company, black-box from our DevOps
- Our TEST environment was able to sustain higher loads without heap issues on ES

*Could it be that the PROD cluster is not properly setup?  
— Anonymous*

# Shift focus: Servers

Metrics were shared from the low-level servers

- Two data-centers, with compute and storage
- ES nodes used storage in the “other” Data-Centers
- ES cluster shared resources with high-demanding applications
  - Which lead to I/O bottlenecks



Fixing the storage config solved 90% of the ES heap cases

# Shift focus: Servers

## Learnings

- Check your underlying infrastructure
- Follow the ES Cluster recommendations!

*I added some highlighting  
to the ES documentation*

*Directly-attached storage performs better than remote storage  
[...] With **careful tuning** it is **sometimes**  
possible to achieve **acceptable performance**  
using remote storage too.*


[https://www.elastic.co/guide/en/elasticsearch/reference/current/tune-for-search-speed.html#\\_use\\_faster\\_hardware\\_2](https://www.elastic.co/guide/en/elasticsearch/reference/current/tune-for-search-speed.html#_use_faster_hardware_2)

*Make sure to use SSDs and  
local storage for ES, avoid NAS!  
— Emanuele Panzeri, 2018*

With age,  
comes with wisdom  
— *Sensei Wu*



# Learnings

- On fire event: keep calm!   
*Check the logs, metrics and monitoring systems*
- Profile your application!  
*Some big optimizations can be quick to implement*
- Trust your Cloud systems  
*Just not always 100%*

# Thank You

